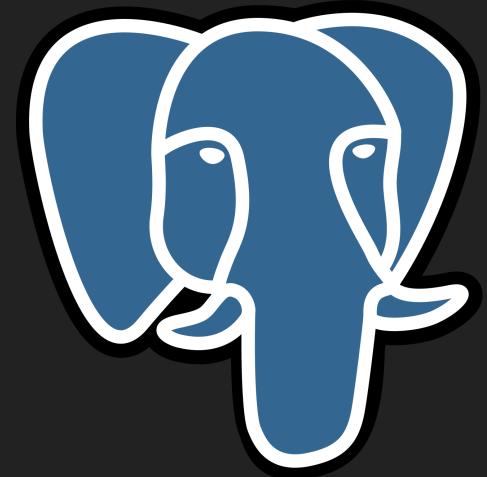


# Puny to Powerful PostgreSQL Rails Apps



Andrew Atkinson

#PgPunyPowerful





# Objectives

- Review 5 types of database scaling challenges and solutions
- Consider integrating the tools and techniques presented into your applications

Safe Migrations on Large, Busy Databases ✨

Maximizing Performance Of Database Connections ⚡

Maintaining Fast SQL Queries ⚡

High Impact Database Maintenance 🔨

Implementing Replication and Partitioning 🐑

# Fountain



- High Volume Hiring Platform
- Remote-first
- RailsConf Sponsor



[fountain.com/careers](https://fountain.com/careers)

[bit.ly/PgPunyPowerful](https://bit.ly/PgPunyPowerful)

Fountain

Product

Solutions

Integrations

Blog

Learn

Company

Sign In

Request Demo

## Find and hire the right people, faster

Exceed your recruiting and hiring goals with the smart, fast, and seamless automation of Fountain.

Request Demo



STITCH FIX

CHIPOTLE

fetch

GORILLAS

gopuff

## Hiring simplified

Fountain combines the power of automation with a data-driven applicant tracking system, so that you can exceed your recruiting and hiring goals.

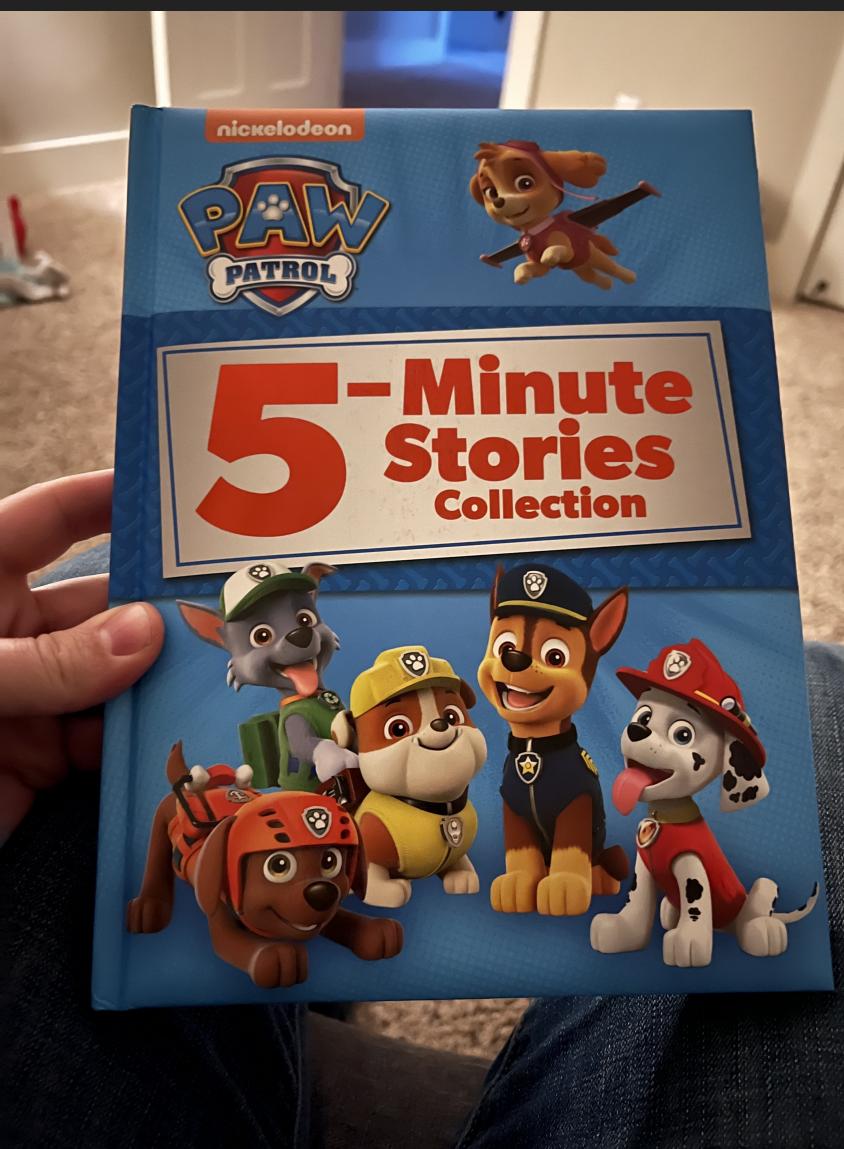
Slide 5 of 69

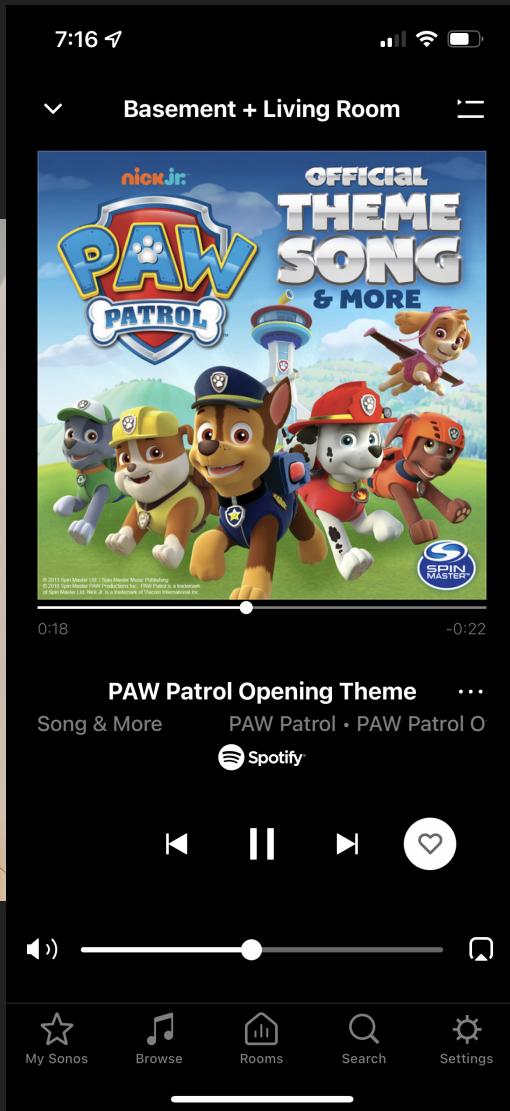


# About Me

- Staff Software Engineer, Fountain
- Work With Rails, PostgreSQL, Elasticsearch
- Dad of 2







[bit.ly/PgPunyPowerful](http://bit.ly/PgPunyPowerful)

Slide 8 of 69

# Paw Patrol 2

## The SQL

[bit.ly/PgPunyPowerful](http://bit.ly/PgPunyPowerful)

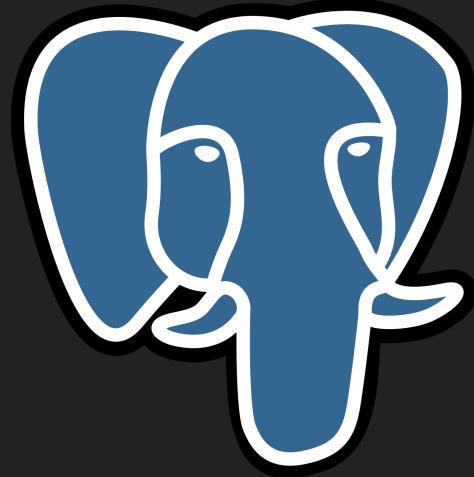
Slide 9 of 69



# Plot

- 5 Years Later
- Mayor Humdinger gone, pups are safe
- Rescues down, Paw Patrol pivoted to startups
- Paw Patrol built Pup Tracker to track pups, keep them safe

# Scaling Pup Tracker



- Pup Tracker built with Rails and PostgreSQL
- Used around the world
- 5 pups share their challenges and solutions





[bit.ly/PgPunyPowerful](https://bit.ly/PgPunyPowerful)

Slide 12 of 69

# Liberty ✨

Newest Paw Patrol  
member, eager to  
make an impact

## Database Migrations

[bit.ly/PgPunyPowerful](https://bit.ly/PgPunyPowerful)



Slide 13 of 69

# Migrations on Large, Busy Tables ✨

As tables grew to millions of rows, making structure changes could take a long time and cause application errors



# Check Constraints: Potentially Unsafe

```
class AddCheckConstraint < ActiveRecord::Migration[7.0]
  def change
    add_check_constraint :pups, "birth_date > '1990-01-01'",
      name: "birth_date_check"
  end
end
```

Adding a check constraint blocks reads and writes in Postgres

Source: **Strong Migrations**

# Check Constraint: Safe Variation

Add the check constraint without validating existing rows

Constraint added as NOT VALID

```
class AddCheckConstraint < ActiveRecord::Migration[7.0]
  def change
    add_check_constraint :pups, "birth_date > '1990-01-01'",
      name: "birth_date_check",
      validate: false
  end
end
```

# Validate Separately

```
class ValidateCheckConstraint < ActiveRecord::Migration[7.0]
  def change
    validate_check_constraint :pups, name: "birth_date_check"
  end
end
```

Less disruptive, but consider running off-peak or in downtime.

| Validation acquires only a SHARE UPDATE EXCLUSIVE lock

Source: Understanding Check Constraints

# Data Backfill: Potentially Unsafe

```
class AddCityIdToLocations < ActiveRecord::Migration[7.0]
  def change
    add_column :locations, :city_id, :integer
    Location.update_all(city_id: 1)
  end
end
```

backfilling in the same transaction that alters a table keeps the table locked for the duration of the backfill.

# Data Backfill: Safe Variation

There are three keys to backfilling safely: batching, throttling, and running it outside a transaction.

Source: [Strong Migrations](#)

# Safe Variation. Migration 1: Add Column.

```
class AddLocationsCityId < ActiveRecord::Migration[7.0]
  def up
    add_column :locations, :city_id, :integer
      # Application sets `city_id` values for new rows
    end
  end
```

# Safe Variation. Second Migration: Backfill

```
class BackfillCityId < ActiveRecord::Migration[7.0]
  disable_ddl_transaction!

  def up
    Location.unscoped.in_batches do |relation|
      relation.update_all(city_id: 1)
      sleep(0.01) # throttle
    end
  end
end
```

# Migrations Good Practices ✨

- Studied potentially unsafe migrations from  
**Strong Migrations**
- Used safe variations
- Modified new rows and modified existing rows  
as separate operations



# Marshall



Emergency responder,  
extinguishes literal fires,  
extinguishes database fires

## Database Connections



# Database Connections



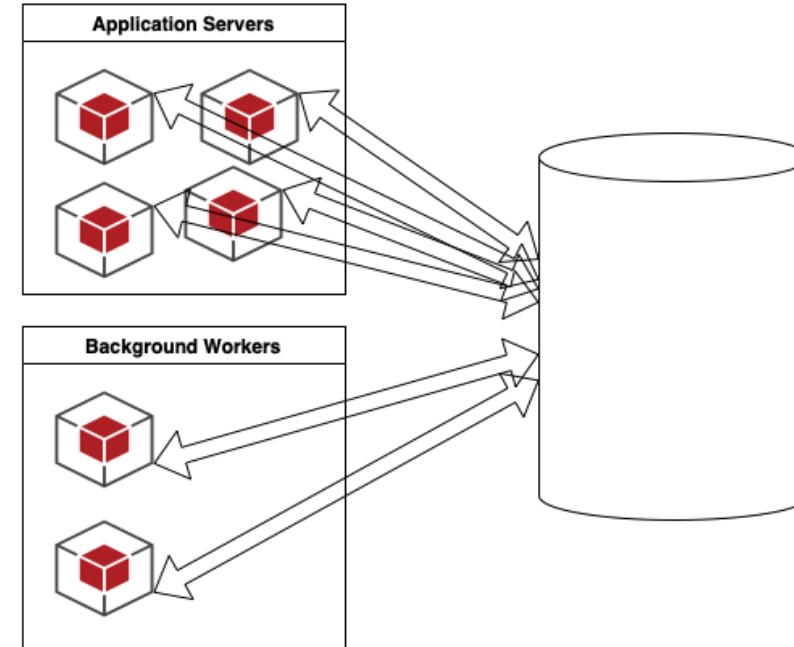
- Marshall learned database connections are finite and may be used up, blocking queries and causing downtime
- Marshall learned that app servers like Puma use connections
- Background workers like Sidekiq use connections



# Active Record Connection Pool

Can set limits

As demand for connections increases, it will create them until it reaches the connection pool limit.



# Active Record Connection Pool and Sidekiq Concurrency



```
# config/database.yml
production:
  pool: 5
```

```
# config/sidekiq.yml
production:
  :concurrency: 5
```

# Resource Consumption



- Each PostgreSQL client connection forks a new process and uses some of the available memory.
- Database server memory is finite
- Besides client connections, memory is also used by background processes and idle connections



# Scaling Options

Marshall needed more memory or more efficient usage of limited connections

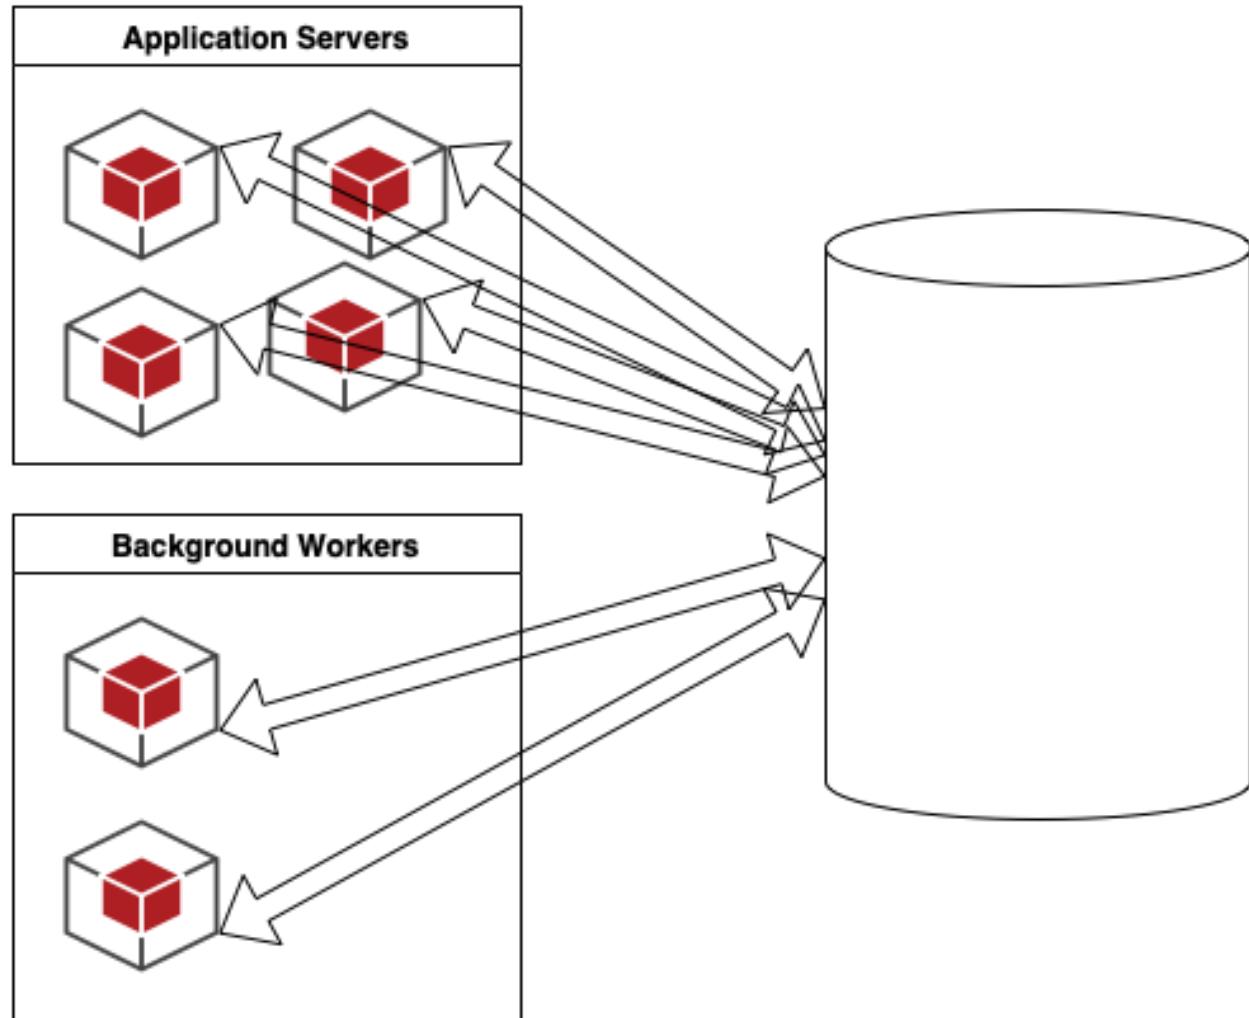
Option #1: Scale database server vertically, gain memory

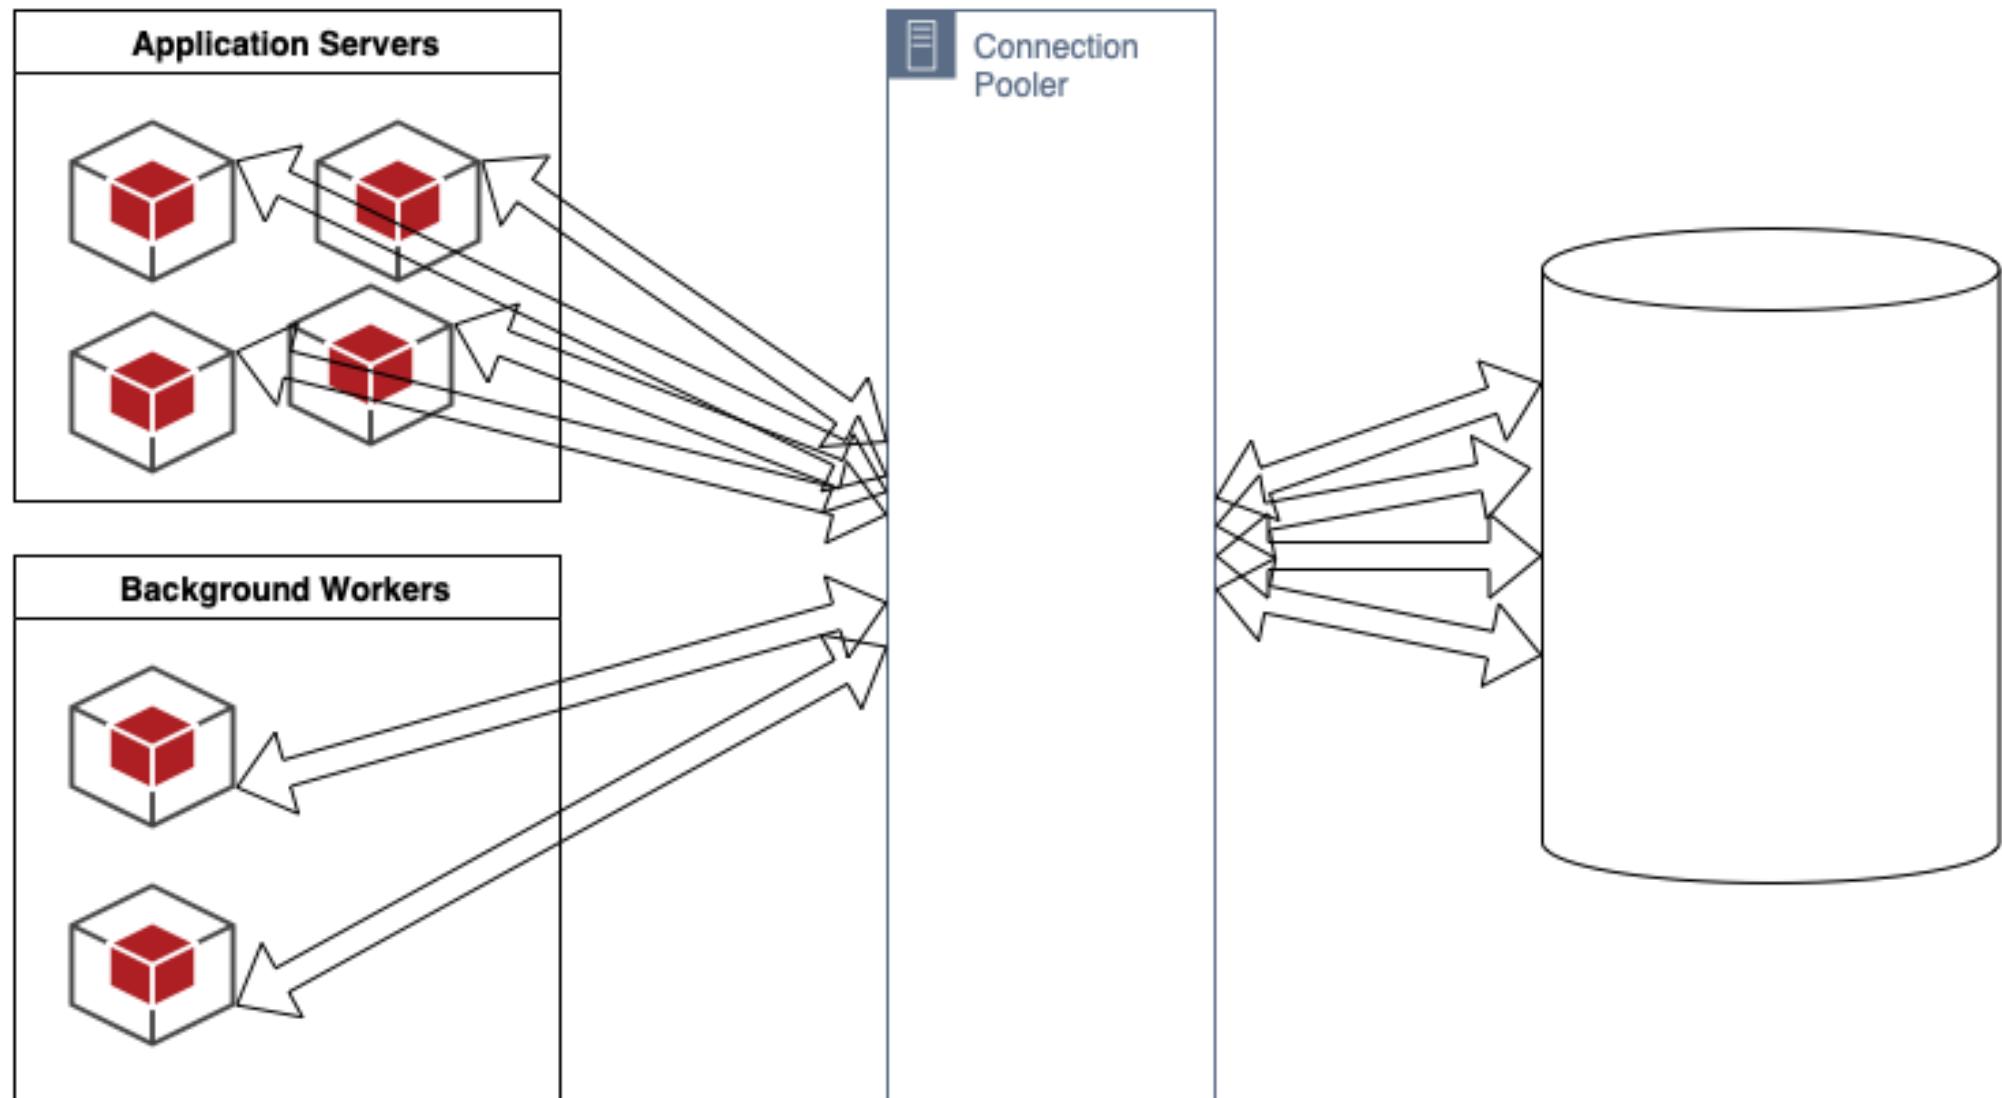
Option #2: Use existing available connections more efficiently

# Connection Poolers



- Connection Poolers are additional software, dedicated to efficient connection usage
- A Connection Pooler does not replace the Active Record Connection Pool
- A Pooler eliminates overhead in establishing new connections
- Poolers offer additional features like pooling modes





# Marshall Implemented Connection Pooling



Marshall installed and configured PgBouncer for Pup Tracker

```
[pgbouncer]
listen_port = 6432
listen_addr = localhost
auth_type = md5
auth_file = /usr/local/etc/userlist.txt
logfile = pgbouncer.log
pidfile = pgbouncer.pid
admin_users = app_user
```

**Connection Pooling improved  
performance without the need to scale  
vertically**



# Chase ⚡

Emergency responder,  
first on the scene

## Fast SQL Queries



# Chase Investigates Slow Queries ⚡

Chase starts by looking at the query plan for a slow query

Example:

```
select * from locations where id = 1;
```

Chase uses `explain` and `explain (analyze)` to explore the query plan and execute the query.



# Query Plan ⚡ (no indexes)

```
select * from locations where id = 1;
```

```
# explain (analyze) select * from locations where id =1;
      QUERY PLAN
-----
Gather  (cost=1000.00..126613.85 rows=1 width=28)
  (actual time=0.474..332.856 rows=1 loops=1)
    Workers Planned: 2
    Workers Launched: 2
    -> Parallel Seq Scan on locations (cost=0.00..125613.75 rows=1 width=28)
        (actual time=215.987..326.496 rows=0 loops=3)
          Filter: (id = 1)
          Rows Removed by Filter: 333333
Planning Time: 0.228 ms
Execution Time: 332.898 ms
```

# Query Plan ⚡ (with primary key constraint/index)

```
alter table locations add primary key (id);
\l locations
Indexes:
"locations_pkey" PRIMARY KEY, btree (id)
```

```
# explain (analyze) select * from locations where id =1;
QUERY PLAN
-----
Index Scan using locations_pkey on locations
(cost=0.43..8.45 rows=1 width=28) (actual time=0.070..0.072 rows=1 loops=1)
  Index Cond: (id = 1)
Planning Time: 0.675 ms
Execution Time: 0.111 ms
```

# Chase Collects Query Stats ⚡

- Which queries are the most impactful? (most resource intensive)
- Need data. Enabled `pg_stat_statements`
- Query normalization (removing parameters):

```
select * from locations where id = ?;
```
- Total Calls, Average Time





A performance  
dashboard for Postgres

Source: PgHero

- Slow Queries
- Unused and Invalid Indexes
- High Connections

The screenshot shows the PgHero dashboard interface. On the left, a sidebar lists navigation options: Overview, Queries (which is selected), Space, Connections, Live Queries, Maintenance, Explain, and Tune. The main area is titled 'Queries' and displays three log entries:

Total Time	Average Time	Calls	User
0 min 28%	11 ms	12	postgres
0 min 28%	10 ms	13	postgres
0 min 10%	46 ms	1	postgres

Each entry includes the SQL query text:

```
SELECT CONCAT(d.first_name, $1, d.last_name) AS driver_name,  
AVG(t.rating) AS avg_rating,  
COUNT(t.rating) AS trip_count  
FROM trips t  
JOIN users d ON t.driver_id = d.id  
GROUP BY t.driver_id, d.first_name, d.last_name  
ORDER BY COUNT(t.rating) DESC
```

```
SELECT schemaname AS schema, t.relname AS table, ix.relname AS name,  
regexp_replace(pg_get_indexdef(i.indexrelid), $1, $2) AS columns,  
regexp_replace(pg_get_indexdef(i.indexrelid), $3, $4) AS using,  
indisunique AS unique, indisprimary AS primary, indisvalid AS valid,  
indexprs::text, indpred::text, pg_get_indexdef(i.indexrelid) AS definition  
FROM pg_index i INNER JOIN pg_class t ON t.oid = i.indrelid INNER JOIN  
pg_class ix ON ix.oid = i.indexrelid LEFT JOIN pg_stat_user_indexes ui ON  
ui.indexrelid = i.indexrelid WHERE schemaname IS NOT NULL ORDER BY 1, 2  
/*pghero*/
```

```
CREATE EXTENSION IF NOT EXISTS pg_stat_statements /*pghero*/
```

# Query Investigation Process ⚡

Get query plan via `EXPLAIN`, make database object changes

OR

- Using **Marginalia** annotations (in Rails 7 🎉), find source location  
`/*controller:locations,action:index*/`
- Make query more selective (add scopes) or rewrite query



# Improving Query Performance ⚡

- Ensured `ANALYZE` ran for table or ran manually (`psql`)
- Added missing indexes (replacing `sequential scan` with `index scan`)
- Rebuilt indexes that were heavily used and had high bloat (`psql`)
- Used performance analysis tools like `pgMustard` and `pganalyze` for insights



# Rubble



Likes construction and  
maintenance

Database  
Maintenance



# Database Maintenance



- `vacuum` ("garbage collect") and Autovacuum
- Updated stats via `analyze` (important for queries)
- Rebuilding Bloated Indexes
- Removing Unused Indexes



# Manually Rebuild Indexes



- Rubble used `REINDEX CONCURRENTLY` on  
`PG >= 12`
- `pg_repack PG <= 11` is another option

## Example

```
# reindex index concurrently locations_rating_idx;  
REINDEX  
Time: 50.108 ms
```



# Automating Maintenance



- Rubble used `pg_cron` extension
- `pg_cron` is cron for database tasks ( PG >= 10 )
- Use the `cron.schedule()` function

```
--- Vacuum every day at 10:00am (GMT) / 3 AM PT  
SELECT cron.schedule('0 10 * * *', 'VACUUM locations');
```

# Database Maintenance

## Recap

- Made sure `vacuum` and `analyze` ran regularly
- Removed unused indexes
- Automated some maintenance tasks



# Skye



Aerial response, takes  
in big picture.

## Multiple Databases

## Replication and Partitioning



# Write and Read Workloads

Reads: **select**

Writes: **insert, update, delete**

- Pup Tracker had 10x more reads than writes
- All writes and reads were on a single database



# Implementing Replication for Read Separation



- Skye set up a read replica to be used for the read workload
- Skye implemented **PostgreSQL streaming replication** (`PG >= 9.0`) (outside Rails)
- Skye used Rails' **Multiple Databases** (`>= 6.0`) to create a primary and replica configuration (inside Rails)



# Database Configuration



config/database.yml

```
production:  
  primary:  
    database: pup_tracker_production  
    adapter: postgresql # obviously!  
  replica:  
    database: pup_tracker_production  
    replica: true
```



Rails will not run migrations against replica

# Application Record



app/models/application\_record.rb

Set the `reading` role

```
class ApplicationRecord < ActiveRecord::Base
  self.abstract_class = true

  connects_to database: {
    writing: :primary,
    reading: :replica
  }
end
```



# Manual Query Relocation



```
ActiveRecord::Base.connected_to(role: :reading) do
  # do a read query here!!
  # e.g. Pup.find_by(name: "Skye").locations.limit(5)
end
```

Consider replication lag

# Automatic Query Relocation



Rails (`>= 6.0`) Automatic role switching

For a GET or HEAD request the application will read from the replica unless there was a recent write.

# Pup Tracker Data Growth



- Pup Tracker only queries most recent month of locations data, but the database has 100s of millions of rows and years of locations
- Query performance has worsened
- Vacuum takes forever to run

# Skye Researches Partitioning



- Skye implemented PostgreSQL Table Partitioning ( PG >= 10 )
- Skye selected Range partitioning,  
created\_at column partition key, monthly partitions



# Partitioning Benefits



- Partitions that no longer reachable by application can be removed
- Greatly reduces quantity of data needed in database
- Smaller indexes, lower cost of table scans, faster maintenance



# Implementation Challenges



- Migration to partitioned tables requires either some downtime or some dual writing (triggers)
- Need to test queries on partitioned table before rollout
- More operational complexity to manage unneeded and upcoming partitions



# Skye used pgslice and pg\_cron



- pgslice generates SQL to help migrate to a partitioned table
- Partition management can be scheduled using pg\_cron



# Some steps



```
CREATE TABLE "locations_intermediate" (LIKE "locations" ...)  
PARTITION BY RANGE ("created_at");
```

```
CREATE TABLE "locations_202206"  
PARTITION OF "locations_intermediate"  
FOR VALUES FROM ('2022-06-01') TO ('2022-07-01');
```

And more. `pgslice` lists all the commands needed!

# Partition Detachment



Partitions can be detached from parent table

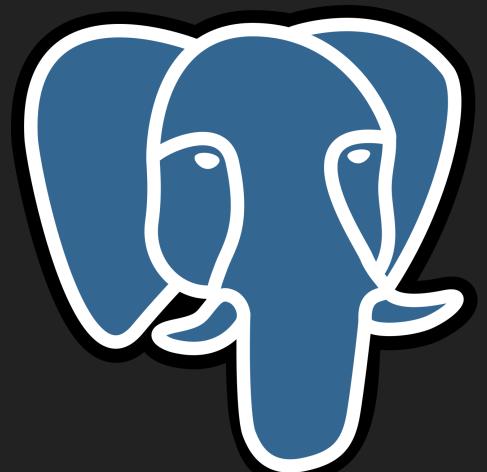
```
ALTER TABLE locations  
DETACH PARTITION locations_20220501 CONCURRENTLY;
```

A detached partition becomes a regular table that can be archived or deleted.



# Summaries

## Database Scaling Challenges and Solutions



# Use Safe Migrations ✨

- Liberty used safe migrations suggested by Strong Migrations
- Liberty separated database changes into two steps, modifying new rows and modifying existing rows



# Use Connection Pooling



- Marshall learned database connections are limited resources
- Marshall implemented a Connection Pooler to improve performance and scalability for Pup Tracker



# Find and Fix Slow Queries ⚡

- Chase set up PgHero as a performance dashboard
- Chase identified and analyzed the most impactful slow queries
- Chase added indexes to improve queries, or sometimes rewrote them



# Database Maintenance



- Rubble monitored background `vacuum` and `analyze` operations, sometimes performing them manually
- Rubble performed manual index rebuilds, sometimes scheduling them with `pg_cron`



# Replication and Partitioning



- Skye created and configured replication between multiple databases
- Skye used the read replica to separate write and read workloads
- Skye implemented range partitioning to cycle out old data and improve performance



# A Special Thank You To Reviewers



@be\_haki

@jmcharnes

@ryanbooz

@davidrowley\_pg

@andrewkane

@andygauge

@dwfrank

@lukasfittl

@kevin\_j\_m

@crash\_tech

👋 Thanks!

#PgPunyPowerful



[bit.ly/PgPunyPowerful](https://bit.ly/PgPunyPowerful)



Slide 68 of 69

- Change new rows and existing rows separately
- Use a connection pooler like PgBouncer
- Fix the most resource intensive slow queries
- Rebuild bloated indexes
- Writes on primary, reads on replica
- Partition very large tables